

NYAY.AI

न्याय

India's First Citizen-Facing Legal AI

Empowering 100 Million Indians to Understand and Exercise Their Legal Rights

50–100M Indians with active legal issues	₹0 access to basic legal information	₹500 Cr+ ARR potential in 5 years
--	--	---

Confidential — For Discussion Purposes Only | May 2026

Table of Contents

Table of Contents2

1. Executive Summary4

2. Understanding the Problem — India's Legal Access Crisis5

 2.1 The Scale of Legal Exclusion5

 2.2 The Five Core Problems We Solve5

 Problem 1: Information Asymmetry5

 Problem 2: Language Barrier5

 Problem 3: Geographic Concentration6

 Problem 4: Cost Barrier6

 Problem 5: Procedural Complexity6

3. Market Analysis — The Size of the Opportunity7

 3.1 Total Addressable Market7

 3.2 Why This Market Is Unserved — Competitive Gap7

 3.3 Reference Markets — Proof It Works8

4. Our Solution — Nyay.AI9

 4.1 Product Vision9

 4.2 Core Product Features9

 Feature 1: Legal Rights Navigator9

 Feature 2: Document Generator9

 Feature 3: Government Scheme Eligibility Engine10

 Feature 4: Lawyer Connect Marketplace10

 Feature 5: Case Status Tracker10

5. Data Sources — Building the Knowledge Base11

 5.1 Primary Legal Corpus11

 Layer 1: Legislation and Statutes11

 Layer 2: Case Law11

 Layer 3: Regulatory Circulars and Guidelines12

 Layer 4: Government Scheme Data12

 Layer 5: Consumer and Citizen-Specific Resources13

 5.2 Data Volume Estimate13

6. The RAG Architecture — How We Build It15

 6.1 Architecture Overview15

 6.2 Data Ingestion Pipeline15

 Step 1: Document Collection15

 Step 2: Document Parsing15

 Step 3: Intelligent Chunking15

 Step 4: Metadata Enrichment16

- 6.3 The Vector Database — Why Hybrid Search is Non-Negotiable 16
- 6.4 LLM Selection and Anti-Hallucination Design 16
 - Primary LLM: Claude claude-sonnet-4 (Anthropic API)..... 16
 - Secondary LLM: Sarvam-2B / Sarvam-M (Indic language queries) 17
 - Anti-Hallucination System Prompt (Non-Negotiable Rules) 17
- 6.5 Confidence Scoring and Escalation Logic..... 17
- 7. Technology Stack — Full Implementation 18
 - 7.1 Stack Overview 18
 - 7.2 Infrastructure Cost Model..... 19
- 8. Build Roadmap — From Zero to Launch 20
 - Phase 1: Foundation (Weeks 1–4)..... 20
 - Phase 2: RAG Chain and API (Weeks 5–8)..... 20
 - Phase 3: WhatsApp + MVP Launch (Weeks 9–12) 20
 - Phase 4: Scale and Monetise (Months 4–12) 21
- 9. Revenue Model and Financial Projections 22
 - 9.1 Revenue Streams 22
 - 9.2 Financial Projections 22
- 10. Go-To-Market Strategy 24
 - 10.1 Distribution Advantage — WhatsApp First..... 24
 - 10.2 Go-To-Market Phases..... 24
 - Phase 1: Community Seeding (Month 1–3)..... 24
 - Phase 2: Earned Media (Month 3–6)..... 24
 - Phase 3: Paid + Partnership Scale (Month 6–12) 24
- 11. Risk Analysis and Mitigation 25
- 12. Team and Execution Requirements 26
 - 12.1 Founding Team Requirements..... 26
 - 12.2 First Hires (Post-Seed)..... 26
 - 12.3 Funding Requirements 26
- 13. Conclusion — Why Now, Why This 28

1. Executive Summary

The Opportunity in One Paragraph

India has 2 million lawyers serving a population of 1.4 billion people. Behind each of those lawyers are millions of ordinary citizens — farmers, tenants, employees, small business owners — who face legal problems every day but have no affordable, accessible, or language-appropriate way to understand their rights or navigate the system. Manupatra, SCC Online, and every legal tech platform built in India in the past 25 years has served the lawyer. Nobody has served the people the lawyers represent. Nyay.AI changes that.

Nyay.AI is a citizen-facing legal AI platform that uses Retrieval-Augmented Generation (RAG) built on India's publicly available legal corpus to give ordinary Indians instant, accurate, multilingual access to their legal rights — delivered via WhatsApp, voice, and a simple web interface. We are not a law firm. We do not give legal advice. We give legal navigation — the difference between a person knowing they have rights and actually being able to exercise them.

Metric	Detail
Product	WhatsApp bot + web app with RAG-powered legal Q&A, rights navigator, and document generator
Primary User	Any Indian citizen with a legal problem — in English, Hindi, or 8 other regional languages
Core Technology	Hybrid BM25 + Vector RAG on 1M+ Indian legal documents, Sarvam AI for Indic languages
Revenue Model	Freemium consumer + lawyer lead marketplace + document generation fees + B2G contracts
Year 1 Target	500,000 free users, 25,000 paid — ₹12 Cr ARR
Year 3 Target	10M users, 500K paid — ₹200 Cr ARR
Primary Moat	Distribution (WhatsApp), language depth, trust built on public legal data, network effects

2. Understanding the Problem — India's Legal Access Crisis

2.1 The Scale of Legal Exclusion

India is simultaneously one of the world's largest democracies and one of its most legally inaccessible countries for ordinary citizens. The Constitution guarantees rights. The legislature passes laws. The courts issue judgments. But the gap between legal rights existing on paper and ordinary Indians being able to exercise them is enormous, structural, and largely ignored by technology.

The Gap	The Reality
50–100 million active legal disputes	Most never reach any formal resolution mechanism
4.7 crore cases pending in courts	Average pendency: 3–5 years at district level
2 million registered advocates	Concentrated in 8–10 major cities; unavailable to 700M rural/semi-urban Indians
Average lawyer consultation fee: ₹2,000–₹5,000	Per capita income in rural India: ~₹1.2 lakh/year — lawyers are unaffordable
22 official languages	Legal system operates almost entirely in English — inaccessible to 80% of the population
1,000+ government schemes	60% of eligible citizens do not know they qualify

2.2 The Five Core Problems We Solve

Problem 1: Information Asymmetry

The most fundamental injustice in India's legal system is information asymmetry. A landlord knows tenant rights better than the tenant. An employer knows labour law better than the worker. A government officer knows citizens' entitlements better than the citizen. This isn't because the information is secret — it's all in public law — but because accessing it requires legal literacy that most Indians simply don't have. A farmer in Vidarbha whose land is being acquired under the Land Acquisition Act doesn't know he's entitled to market value plus a 100% solatium. He accepts whatever is offered. The law protects him. Nobody told him.

Problem 2: Language Barrier

India's legal system operates in English — a language that less than 20% of Indians speak with any proficiency. High Court judgments are in English. The Income Tax Act is in English. RTI applications are ideally filed in English. Consumer forum complaints are more effective in English. This creates a structural barrier where the quality of justice a citizen can access is directly correlated with their English proficiency and educational background. A Tamil-speaking auto driver in Coimbatore has a legitimate consumer complaint but cannot navigate a system that communicates entirely in a language he doesn't understand.

Problem 3: Geographic Concentration

India has roughly 2 million lawyers for 1.4 billion people — a ratio of 1:700 on paper. In practice, the distribution is wildly skewed. Over 60% of practicing advocates are concentrated in Delhi, Mumbai, Bangalore, Chennai, Kolkata, and Hyderabad. A person in a Tier-3 city or village who needs legal help faces a genuine shortage of affordable, qualified lawyers. The lawyers available locally may be competent for routine matters but often lack exposure to complex statutory law, regulatory compliance, or constitutional rights.

Problem 4: Cost Barrier

A straightforward legal consultation in a metro city costs ₹2,000–₹5,000 for 30 minutes. A demand notice drafted by a lawyer costs ₹3,000–₹10,000. A consumer forum complaint with legal assistance costs ₹5,000–₹15,000. For a worker earning ₹15,000/month, the cost of asserting a legitimate legal right through the formal system can easily equal one month's salary — before the case even begins. The result: most people don't pursue legal remedies even when they clearly have them. They absorb the injustice because the alternative is too expensive.

Problem 5: Procedural Complexity

Even when a person knows their rights and can afford a lawyer, the procedural complexity of India's legal system is genuinely daunting. Which forum has jurisdiction? What documents are needed? What is the limitation period? Which form must be filed? In which language? To which officer? India's legal machinery was designed by and for trained professionals. For a first-generation litigant with no prior exposure, each of these questions is a potential fatal error. Many legitimate cases are thrown out on procedural grounds before they are ever heard on merit.

The Core Insight

Every existing legal tech platform in India — Manupatra, SCC Online, CaseMine, LegitQuest — was built to make lawyers more efficient. None of them was built to make lawyers unnecessary for the 80% of legal problems that don't actually require a lawyer. The market they are all fighting over is ₹200–500 Cr. The market that nobody is serving is ₹5,000–10,000 Cr.

3. Market Analysis — The Size of the Opportunity

3.1 Total Addressable Market

The TAM for citizen-facing legal AI in India can be calculated from multiple angles, all of which converge on a very large number.

Segment	Size	Revenue Potential
Consumer subscriptions (₹99–499/mo)	50M addressable users, 2% paid = 1M subscribers	₹600–2,400 Cr/year
Lawyer lead marketplace	20L advocates needing clients; 1M+ annual referrals @ ₹1,000/lead avg	₹100–500 Cr/year
Document generation (₹49–499/doc)	Crores of legal notices, RTIs, complaints filed annually	₹200–1,000 Cr/year
B2G (state/central govt contracts)	Legal aid scheme digitisation, scheme awareness	₹50–200 Cr/year
B2B (corporates, NBFCs, banks)	Employee legal helplines, customer complaint resolution	₹100–300 Cr/year
TOTAL TAM		₹1,000–4,000 Cr/year by 2030

3.2 Why This Market Is Unserved — Competitive Gap

Indian Kanoon	Free search, no AI, no guidance, no documents, no Hindi — pure case law database for researchers
Manupatra / SCC Online	Sophisticated tools for lawyers; subscription ₹6,500–₹60,000/year; English-only; zero citizen orientation
CaseMine / LegitQuest	AI-powered research for advocates; still entirely B2B; no consumer product
Vakil Search / LawRato	Lawyer directories; no legal information; pure marketplace with no intelligence layer
Legal Services Authority	Government NALSA free legal aid; massively under-resourced; no technology layer; requires physical presence
No one	AI-powered, WhatsApp-delivered, multilingual, citizen-facing legal information and navigation — this is the gap

3.3 Reference Markets — Proof It Works

The citizen-facing legal AI model has been validated in other markets with comparable dynamics:

- DoNotPay (US/UK): Started as a parking ticket appeal bot, grew to 2M+ users handling consumer complaints, subscription cancellations, and landlord disputes. Raised \$27M before going private. India's market is 4x larger with far less competition.
- Rocket Lawyer (US): \$223M in revenue serving non-lawyer legal document generation. Proves massive willingness to pay for affordable legal document services among consumers.
- Kaanoon (India, early stage): Basic legal Q&A startup gaining traction — validates Indian consumer demand but has no AI layer, no WhatsApp delivery, no regional language support.
- MyAdvo (India): Lawyer referral platform; raised ₹30 Cr; proves Indians will pay for legal navigation help online.

4. Our Solution — Nyay.AI

4.1 Product Vision

"Legal navigation for every Indian, in their language, in their pocket."

Nyay.AI is not a law firm. It does not give legal advice. It gives legal navigation — helping citizens understand what the law says, what their options are, what documents they need, and when they actually need a lawyer. It is the difference between being lost in a foreign city and having Google Maps.

4.2 Core Product Features

Feature 1: Legal Rights Navigator

A conversational interface (WhatsApp + web) where a citizen describes their situation in plain language — in Hindi, English, Tamil, Telugu, Kannada, Marathi, Bengali, Gujarati, Malayalam, or Punjabi — and receives a clear explanation of their legal rights, applicable laws, and available remedies. The system retrieves the relevant statutes, landmark judgments, and regulatory guidelines and synthesises them into plain-language guidance with source citations.

Example interaction:

User (in Hindi): "मेरे मकान मालिक ने बिना नोटिस के बिजली काट दी। क्या यह कानूनी है?"

Nyay.AI: Your landlord cannot disconnect electricity without a court order or proper notice under the Transfer of Property Act 1882 and the Electricity Act 2003 (Section 56). This is a punishable offence. Your immediate options are: (1) File a complaint with the District Consumer Forum under the Consumer Protection Act 2019 (free to file, no lawyer needed for claims under ₹50L), (2) Send a legal notice via our platform (₹99), (3) File a police complaint under IPC Section 340 for wrongful restraint. Would you like us to generate the consumer forum complaint for you now?

Feature 2: Document Generator

Pre-built, legally validated templates for the 50 most common legal documents an Indian citizen needs, with intelligent auto-fill based on the conversation context. No lawyer needed. Documents are drafted, formatted, and downloadable as PDF or Word within minutes.

Document Type	Price	Annual Filing Volume (Est.)
RTI Application	₹49	7 crore+ RTI applications filed annually
Legal Notice (Demand)	₹99	5 crore+ demand notices sent annually
Consumer Forum Complaint	₹149	60 lakh+ consumer complaints annually
Labour / Wage Complaint	₹149	20 lakh+ labour complaints annually

Document Type	Price	Annual Filing Volume (Est.)
Police Complaint Draft	₹49	Crores of FIR / complaint drafts needed
Tenant Eviction Notice	₹149	Massive volume in metros
RERA Homebuyer Complaint	₹199	40 lakh+ RERA complaints pending
Cheque Bounce (Sec 138) Notice	₹199	40 lakh+ cheque bounce cases/year

Feature 3: Government Scheme Eligibility Engine

A conversational eligibility checker covering 500+ central and state government schemes. User inputs their state, income, occupation, caste category, and family situation; the system retrieves applicable schemes, explains eligibility criteria, and generates the application checklist with form links. This feature alone serves a massive underserved population — an estimated 60% of eligible Indians do not claim benefits they are legally entitled to.

Feature 4: Lawyer Connect Marketplace

When the user's situation requires actual legal representation, Nyay.AI connects them with a verified, rated local advocate. Unlike generic directories, the referral is contextual — the lawyer receives a structured brief of the user's situation, applicable laws, and documents already prepared by the platform. This dramatically reduces the lawyer's onboarding time and makes the referral 10x more valuable than a cold directory listing. Lawyers pay per qualified referral.

Feature 5: Case Status Tracker

Integration with eCourts API to allow citizens to track their own case status, next hearing dates, and orders — without needing to call their lawyer or visit the court. Delivered via WhatsApp notifications. Lawyers can subscribe to premium tracking alerts for their entire case portfolio.

5. Data Sources — Building the Knowledge Base

The foundation of any RAG system is the quality, coverage, and freshness of its underlying data. Nyay.AI is built entirely on publicly available Indian legal data — which means no data licensing costs, no vendor dependency, and no proprietary lock-in. All sources below are legally and freely accessible.

5.1 Primary Legal Corpus

Layer 1: Legislation and Statutes

Source	Description & Access Method
India Code (indiacode.nic.in)	All central acts, amendments, and bare acts — 1,500+ statutes available as free XML and PDF downloads. Updated within weeks of presidential assent. This is the authoritative source for all central legislation.
State Legislature Portals	Each state has its own portal (e.g., maharashtra.gov.in/legislature, vidhansabha.kar.nic.in). State-specific laws like Rent Control Acts, Shops and Establishments Acts, and Agricultural Land Acts are available as free PDFs.
Ministry/Department Websites	Labour Ministry, MCA, SEBI, RBI, TRAI, IRDAI, Competition Commission — all publish their enabling legislation, rules, and delegated regulations. Systematically scrape each ministry's 'Acts & Rules' section.
New Criminal Laws (2024)	BNS, BNSS, and BSA (replacements for IPC, CrPC, Evidence Act) are on India Code. Critical to index separately given recency — courts are still adapting.

Layer 2: Case Law

Source	Description & Access Method
Supreme Court (sci.gov.in)	All SC judgments from 1950 onwards. Free PDF downloads organised by year and month. ~85,000+ judgments. Build an automated scraper with Scrapy + PyPDF.
eCourts API (ecourts.gov.in)	Official government API for High Court and District Court judgments and case status. Free registration required. Provides structured data including case metadata, parties, judge names, dates — ideal for RAG metadata filtering.
High Court Websites (25 HCs)	Each High Court has its own judgment portal. Priority order: Delhi HC, Bombay HC, Madras HC, Karnataka HC, Calcutta

	HC (highest case volumes and English judgments). Many HCs now also upload in regional languages.
Indian Kanoon (indiankanoon.org)	Aggregates judgments from most courts. Good as a supplementary cross-reference and to fill gaps in official portals. Has a reasonable robots.txt — scrape with care and rate limiting.
ITAT (itat.gov.in)	Income Tax Appellate Tribunal orders — critical for any tax-related guidance. Free public access.
NCLT / NCLAT (nclt.gov.in)	National Company Law Tribunal — company law, insolvency, merger cases. Critical for any business/property related queries.
Consumer Forums (ncdrc.nic.in)	National Consumer Disputes Redressal Commission orders. Highly relevant for a citizen-facing platform since this is where ordinary people most often litigate.

Layer 3: Regulatory Circulars and Guidelines

Source	Description & Access Method
RBI Master Circulars (rbi.org.in)	All FEMA, banking, NBFC, payment system regulations. Updated annually in July. Critical for any banking/loan related queries.
SEBI Circulars (sebi.gov.in)	All capital market regulations, investor protection guidelines. Essential for securities law queries.
GSTN / CBIC (gst.gov.in)	All GST notifications, circulars, AAR advance rulings — over 2,000 documents since 2017.
Labour Ministry Circulars (labour.gov.in)	4 Labour Codes implementation notifications, EPF/ESIC circulars, minimum wage notifications (state-wise).
MCA Circulars (mca.gov.in)	Company law interpretations, LLP rules, director liability guidelines.
IRDAI (irdai.gov.in)	Insurance regulations — critical for claim denial and policy mis-selling queries, a major consumer complaint category.

Layer 4: Government Scheme Data

Source	Description & Access Method
--------	-----------------------------

MyScheme Portal (myscheme.gov.in)	Official government scheme aggregator covering 500+ schemes with eligibility criteria. Has a structured API — the cleanest source for scheme data.
State Government Portals	Each state has 100–300 additional state-specific schemes. Systematically crawl each state's welfare/schemes portal (varies by state — Maharashtra has mahaonline.gov.in, Karnataka has serviceonline.gov.in, etc.)
PM India (pmindia.gov.in)	All central schemes with FAQs, eligibility, and application links.
JAM Trinity Data	Jan Dhan, Aadhaar, Mobile — scheme disbursement and coverage data from DBT.gov.in.

Layer 5: Consumer and Citizen-Specific Resources

Source	Description & Access Method
PIB Fact-Checks (pib.gov.in)	Government fact-checks and press releases — useful for correcting legal misinformation that spreads via WhatsApp.
NALSA Legal Aid Schemes (nalsa.gov.in)	National Legal Services Authority guidelines — who qualifies for free legal aid, how to apply.
Consumer Helplines	National Consumer Helpline (1800-11-4000) FAQs, complaint categories, resolution rates — useful training data.
RTI Guidelines (cic.gov.in)	Central Information Commission orders and RTI Act interpretation — critical since RTI is one of the highest-volume citizen legal actions in India.

5.2 Data Volume Estimate

Data Layer	Est. Documents	Est. Size (GB)	Update Freq.
Central Acts + State Acts	2,500+	8 GB	Monthly
SC Judgments	85,000+	15 GB	Daily
HC Judgments (top 5 HCs)	500,000+	60 GB	Daily
Tribunal Orders (ITAT, NCLT etc.)	200,000+	25 GB	Weekly
Consumer Forum Orders	100,000+	12 GB	Weekly

Data Layer	Est. Documents	Est. Size (GB)	Update Freq.
Regulatory Circulars	10,000+	5 GB	Weekly
Government Schemes	800+	2 GB	Monthly
TOTAL	~900,000 docs	~127 GB raw	Incremental daily

Important Note on Data Costs

All data sources listed above are legally and freely accessible under the National Data Sharing and Accessibility Policy (NDSAP, 2012) and the Right to Information Act. Government data is crown copyright material published for public access. There are zero licensing costs for building on top of this corpus.

6. The RAG Architecture — How We Build It

RAG (Retrieval-Augmented Generation) is the correct architecture for legal AI for a specific reason: it grounds every answer in retrieved, verifiable source documents rather than relying on an LLM's parametric memory. This is non-negotiable in legal contexts where a single hallucinated citation can cause catastrophic harm to a user's case. Every answer Nyay.AI produces is traceable to a specific document, paragraph, and source URL.

6.1 Architecture Overview

Nyay.AI RAG Pipeline — Step by Step

1. USER QUERY → Language detection + translation to English (Sarvam AI / Google Translate)
2. QUERY UNDERSTANDING → LLM extracts: legal domain, jurisdiction, applicable acts, court level, time sensitivity
3. HYBRID RETRIEVAL → BM25 (keyword) + Dense Vector (semantic) search on Qdrant, filtered by metadata
4. RERANKING → Cohere Rerank API or BGE-Reranker-Large sorts retrieved chunks by relevance
5. CONTEXT ASSEMBLY → Top-8 chunks assembled with source metadata into structured context
6. LLM GENERATION → Claude claude-sonnet-4 generates answer grounded strictly in retrieved context
7. CITATION INJECTION → Every factual claim linked to specific source document + paragraph
8. LANGUAGE OUTPUT → Response translated back to user's query language if not English
9. ESCALATION LOGIC → If confidence < threshold, recommend lawyer consultation

6.2 Data Ingestion Pipeline

Step 1: Document Collection

Automated scrapers run on a scheduled basis (daily for case law, weekly for circulars, monthly for legislation updates). Each scraper is purpose-built for its source:

- Scrapy-based spiders for each court portal and government website
- Selenium for JavaScript-heavy portals that block simple HTTP scrapers
- Official APIs where available (eCourts API, MyScheme API, PIB RSS feeds)
- PDF download pipelines with retry logic and deduplication by SHA-256 hash

Step 2: Document Parsing

Raw PDFs and HTML pages are transformed into clean text using a parsing pipeline:

- Unstructured.io for multi-format parsing (PDF, DOCX, HTML, scanned images)
- Tesseract OCR for scanned legacy judgments (pre-2000 SC decisions)
- Custom post-processing to remove headers, footers, page numbers, and watermarks from court PDFs
- Language detection using langdetect; flag non-English documents for translation

Step 3: Intelligent Chunking

Standard chunking splits documents arbitrarily — catastrophic for legal text where a sentence referencing Section 138 is meaningless without the preceding statutory context. Nyay.AI uses legal-aware chunking:

- Target chunk size: 600–800 tokens with 15% overlap
- Paragraph-boundary respecting — never split mid-paragraph
- Section-header anchoring — each chunk includes its parent section heading as context prefix
- For judgments: separate chunks for facts, arguments, holding, and ratio decidendi using NLP classification
- For acts: each section is a separate chunk with the act name, section number, and full text

Step 4: Metadata Enrichment

Every chunk is tagged with rich metadata that enables precise filtered retrieval:

Example metadata object per chunk:

```
{
  doc_type: 'judgment',      // judgment | act | circular | scheme | guideline
  court: 'Supreme Court',    // SC | HC-Delhi | HC-Bombay | ITAT | NCLT | NCDRC
  jurisdiction: 'national',  // national | state:Maharashtra | state:Karnataka
  domain: 'consumer',       // consumer | labour | property | criminal | family | tax | company
  year: 2024,
  citation: 'AIR 2024 SC 1234',
  acts_cited: ['Consumer Protection Act 2019 s.35', 'CPC Order 7 Rule 11'],
  is_binding: true,          // SC binding | HC persuasive | tribunal non-binding
  source_url: 'https://sci.gov.in/...',
  language: 'en',           // original language of document
  chunk_role: 'holding'     // facts | arguments | holding | statute_text | scheme_eligibility
}
```

6.3 The Vector Database — Why Hybrid Search is Non-Negotiable

Pure semantic (vector) search fails in legal contexts. A query for 'Section 138 NI Act 60 days notice' relies on exact statutory references that embedding models may treat as semantically equivalent to unrelated concepts. Conversely, pure keyword (BM25) search misses conceptually related cases that use different terminology.

Nyay.AI uses Qdrant's native hybrid search combining:

- Dense vectors (multilingual-e5-large, 1024 dimensions) for semantic similarity — free, runs on-premise, supports 10 Indian languages
- Sparse vectors (BM25 via Qdrant's sparse vector support) for exact keyword matching of section numbers, act names, citation strings
- Metadata filtering applied pre-retrieval to limit search scope (e.g., filter to consumer law + Maharashtra + post-2020)
- Reciprocal Rank Fusion (RRF) to merge dense and sparse results into a single ranked list

6.4 LLM Selection and Anti-Hallucination Design

Primary LLM: Claude claude-sonnet-4 (Anthropic API)

Claude is the correct choice for this application for three specific reasons: (1) it has a 200K token context window, allowing us to pass full judgment texts as context rather than fragments; (2) it follows complex system prompt instructions reliably, critical for enforcing citation-only output rules; and (3) it has strong instruction-following on safety boundaries, reducing risk of producing responses that veer into actual legal advice.

Secondary LLM: Sarvam-2B / Sarvam-M (Indic language queries)

For queries in Hindi, Tamil, Telugu, Kannada, Marathi, Bengali, Gujarati, Malayalam, and Punjabi, we route through Sarvam AI's API, which is specifically trained on Indic language legal text. This is superior to Google Translate + English LLM for nuanced regional queries.

Anti-Hallucination System Prompt (Non-Negotiable Rules)

SYSTEM PROMPT — Nyay.AI Legal Navigator

You are Nyay.AI, a legal navigation assistant for Indian citizens. You MUST follow these rules without exception:

RULE 1: Only cite cases, sections, or circulars that appear VERBATIM in the RETRIEVED CONTEXT below. If it is not in the context, do not cite it.

RULE 2: Never say 'you should' or 'I advise you to' — you provide information, not advice. Use: 'the law provides', 'your options include', 'under Section X, a citizen may'.

RULE 3: If the retrieved context does not adequately answer the question, say so explicitly and recommend consultation with a qualified advocate.

RULE 4: Always end with: 'This is legal information, not legal advice. For your specific situation, consult a qualified advocate.'

RULE 5: Temperature = 0. No creativity in legal facts.

6.5 Confidence Scoring and Escalation Logic

Not all queries can be safely answered by the RAG system. Nyay.AI implements a confidence scoring layer that determines when to answer, when to caveat, and when to escalate to a human lawyer:

Confidence Level	Trigger Condition	System Action
HIGH (>0.85)	Retrieved chunks directly address the query; multiple consistent sources	Generate full answer with citations
MEDIUM (0.6–0.85)	Partial match; some relevant context found; jurisdiction/time uncertainty	Generate answer with explicit caveats; highlight uncertainties
LOW (<0.6)	Query too fact-specific; no relevant context retrieved; criminal charges	Decline to answer; explain limitation; offer lawyer connect
ALWAYS ESCALATE	Criminal charges facing user; custody; child custody; domestic violence	Provide emergency helpline numbers first, then offer lawyer connect

7. Technology Stack — Full Implementation

7.1 Stack Overview

Layer	Technology	Rationale
Data Collection	Scrapy, Playwright, Requests, schedule	Headless scraping + scheduled jobs for all sources
Document Parsing	Unstructured.io, Tesseract OCR, PyMuPDF	Handles PDF, scanned docs, HTML — best Indian court PDF parser
Embeddings	multilingual-e5-large (HuggingFace)	Free, on-premise, excellent Hindi/Tamil/Telugu performance
Vector DB	Qdrant Cloud (ap-south-1 region)	Best OSS vector DB; hybrid search; payload filtering; free tier
BM25 Index	Qdrant sparse vectors / BM42	Native hybrid search without separate Elasticsearch cluster
Reranking	Cohere Rerank API / BGE-Reranker-Large	15-20% accuracy lift on Indian legal queries; critical for precision
Primary LLM	Claude claude-sonnet-4 (Anthropic API)	200K context; strong instruction following; reliable citations
Indic LLM	Sarvam AI API	Best Hindi/Tamil/Telugu/Kannada legal text understanding
RAG Orchestration	LlamaIndex (document pipelines) + LangChain (chains)	LlamaIndex for ingestion, LangChain for query-time orchestration
Backend	FastAPI (Python 3.11)	Async, high performance, native Pydantic validation
Auth	Supabase Auth (JWT + RLS)	Free tier; row-level security for multi-tenant user data
Database	Supabase PostgreSQL + pgvector	User data, conversation history, document metadata, quota tracking
WhatsApp	Meta Cloud API (WhatsApp Business)	Official API; 500M Indian users; free to receive; ₹0.45–0.89/message sent
Frontend	Next.js 15 + Tailwind CSS	Fast, SEO-optimised, server components for low-latency first load

Layer	Technology	Rationale
Payments	Razorpay (subscriptions + one-time)	India-native; UPI support; lowest failure rates for Indian cards
Cloud	AWS ap-south-1 (Mumbai)	Data residency; lowest latency; compliance with Indian data localisation
Monitoring	Langfuse (LLM observability) + Sentry	Token cost per user, query latency, error tracking, per-user analytics
CI/CD	GitHub Actions + Docker + ECS Fargate	Serverless containers; auto-scaling; zero infra management

7.2 Infrastructure Cost Model

Cost control is critical for a consumer-facing product targeting price-sensitive Indian users. The following estimates are based on AWS ap-south-1 pricing and current API rates:

Component	Monthly Cost (MVP)	Scales to
Qdrant Cloud (8GB)	\$199/mo	\$600/mo at 100GB (1M+ docs)
AWS ECS Fargate (2 tasks)	~\$80/mo	\$400/mo at 10k users/day
Supabase Pro	\$25/mo	\$599/mo at scale
Claude API (10K queries/day)	~\$90/mo (\$0.003/query avg)	\$900/mo at 100K queries/day
Sarvam AI API (20% of queries)	~\$30/mo	\$300/mo at scale
WhatsApp Business API	₹0.50/msg sent; free to receive	Scales linearly — price into per-query cost
Cohere Rerank API	~\$40/mo	\$400/mo at scale
Cloudflare (CDN + Pages)	\$20/mo	Flat — unlimited bandwidth
TOTAL MVP	~\$484/mo (₹40,000/mo)	~\$3,400/mo at 10k DAU

8. Build Roadmap — From Zero to Launch

Phase 1: Foundation (Weeks 1–4)

Goal: Working RAG system on core legal corpus. No UI. Validate retrieval quality.

1. Set up Qdrant Cloud cluster in ap-south-1. Create collection 'nyay_legal' with hybrid search configuration (dense + sparse vectors).
2. Write Scrapy spiders for: India Code (all central acts), Supreme Court judgments (2010–2026), eCourts API (HC judgments — Delhi, Bombay, Madras).
3. Build PDF parsing pipeline using Unstructured.io + PyMuPDF. Process and clean ~200,000 initial documents.
4. Run chunking pipeline: 700 token chunks, 15% overlap, section-boundary respecting, metadata enrichment for all chunks.
5. Generate embeddings using multilingual-e5-large. Upsert all chunks to Qdrant with full metadata payload.
6. Build BM42 sparse vector index for the same corpus within Qdrant.
7. Test retrieval quality: run 100 test queries across different legal domains; measure Recall@5 and MRR. Target: Recall@5 > 0.75.
8. Tune RRF weights (alpha parameter balancing dense vs sparse) for optimal hybrid retrieval on Indian legal text.

Phase 2: RAG Chain and API (Weeks 5–8)

Goal: Accurate, citation-grounded answers via REST API. Ready for WhatsApp integration.

9. Build LangChain RAG chain: query understanding → hybrid retrieval → Cohere reranking → Claude generation with anti-hallucination system prompt.
10. Implement query classification: domain detection (consumer/labour/property/criminal/family/tax), jurisdiction detection, urgency detection.
11. Implement confidence scoring system. Wire escalation logic for low-confidence queries and emergency situations.
12. Build FastAPI backend: /query endpoint, /document-generate endpoint, /scheme-check endpoint. Add Supabase JWT auth.
13. Add freemium quota enforcement: 10 free queries/day per phone number; unlimited for paid users.
14. Set up Langfuse for LLM observability: track token usage, cost per query, latency, citation quality.
15. Integrate Sarvam AI for Indic language routing: detect non-English queries → route to Sarvam; translate response back.
16. Build document generation module: 10 document templates (RTI, legal notice, consumer complaint, wage complaint, cheque bounce notice).

Phase 3: WhatsApp + MVP Launch (Weeks 9–12)

Goal: Public WhatsApp bot + web interface live. First 1,000 real users. Revenue starts.

17. Register WhatsApp Business account and get Meta Cloud API access. Set up webhook to FastAPI backend.
18. Build WhatsApp conversation flow: greeting → problem description → domain classification → RAG answer → document offer → lawyer connect.
19. Build minimal Next.js web frontend: landing page, chat interface, document download, Razorpay payment integration.
20. Integrate Razorpay subscriptions: ₹99/mo Basic (unlimited queries), ₹299/mo Pro (queries + documents), ₹999/mo premium lawyer referral priority.
21. Set up automated daily scraper jobs (AWS EventBridge → Lambda → scraper containers) for case law freshness.
22. Deploy to AWS ECS Fargate (ap-south-1). Configure CloudFront CDN for frontend. Set up domain, SSL, monitoring.
23. Soft launch to 100 beta users recruited from bar association networks, law school student groups, and consumer forum communities. Collect feedback.
24. Fix top-10 issues from beta. Improve response quality for top query categories. Launch publicly.

Phase 4: Scale and Monetise (Months 4–12)

- Expand corpus to all 25 High Courts, all major tribunals, consumer forum orders, RERA portals
- Launch lawyer marketplace: onboard 500+ verified advocates in top 20 cities, build rating/review system
- Add voice input (Bhashini API integration for voice-to-text in 10 languages)
- Government scheme eligibility engine covering 500+ central and state schemes
- B2B product: API for NBFCs, banks, legal aid NGOs to embed Nyay.AI in their own interfaces
- Expand document templates to 50 document types covering all major consumer/citizen use cases
- Target first B2G contract: approach NALSA (National Legal Services Authority) or state legal aid board for partnership

9. Revenue Model and Financial Projections

9.1 Revenue Streams

Stream	Price Point	How It Works
Free Tier	₹0	10 queries/day via WhatsApp. Drives acquisition. Monetised through lead gen.
Basic Subscription	₹99/mo	Unlimited queries, Hindi/English. Target: individual citizens with recurring legal needs.
Pro Subscription	₹299/mo	Unlimited queries + document generation (RTI, notices, complaints). Priority responses.
Document Generation (à la carte)	₹49–₹499/doc	One-time purchase. No subscription needed. Massive volume opportunity.
Lawyer Lead Referral	₹500–₹2,500/lead	Contextual referral with pre-built case brief. Premium over generic directories.
B2B API Access	₹15,000–₹1L/mo	NBFC employee legal helplines, bank customer grievance resolution, NGO legal aid.
B2G Contracts	₹10L–₹5Cr/contract	State legal aid digitisation, scheme awareness campaigns, consumer helpline integration.

9.2 Financial Projections

Metric	Year 1	Year 2	Year 3
Total Users (free + paid)	5,00,000	30,00,000	1,00,00,000
Paid Subscribers	10,000	1,50,000	5,00,000
Documents Generated	50,000	5,00,000	25,00,000
Lawyer Referrals	5,000	75,000	3,00,000
Subscription Revenue	₹1.2 Cr	₹18 Cr	₹60 Cr
Document Revenue	₹0.5 Cr	₹5 Cr	₹25 Cr
Lead Marketplace Revenue	₹0.5 Cr	₹7.5 Cr	₹30 Cr

Metric	Year 1	Year 2	Year 3
B2B + B2G Revenue	₹1 Cr	₹10 Cr	₹50 Cr
TOTAL ARR	₹3.2 Cr	₹40.5 Cr	₹165 Cr
Gross Margin	55%	68%	75%
Monthly Burn (operations)	₹8–12 Lakh	₹30–50 Lakh	₹1–2 Cr

10. Go-To-Market Strategy

10.1 Distribution Advantage — WhatsApp First

WhatsApp is the single most important distribution decision in this business. India has 500+ million WhatsApp users — more than the entire population of the United States. For our target users (Tier-2/3 city residents, rural users, first-generation smartphone users), WhatsApp is the internet. They do not install new apps. They do not visit websites. They chat on WhatsApp. Building Nyay.AI as a WhatsApp-first product is not a feature — it is the distribution strategy.

A WhatsApp number is as memorable and shareable as a phone number. When a Nyay.AI user solves their problem and tells their neighbour, they share a phone number — not an app store link. This is viral distribution with zero friction.

10.2 Go-To-Market Phases

Phase 1: Community Seeding (Month 1–3)

- Target consumer forum communities on Facebook and WhatsApp — 1,000+ active groups discussing landlord disputes, salary theft, insurance claim denials
- Partner with 5 consumer law YouTubers who have 100K+ Hindi-language subscribers — demonstration videos of real query resolution
- District Bar Association demos in 5 cities — position as a tool that pre-qualifies clients and saves lawyers time on basic queries
- Legal aid NGOs (iCall, HRLN, Prayas) — offer free API access in exchange for testimonials and user referrals

Phase 2: Earned Media (Month 3–6)

- Target The Hindu, Indian Express, Wire for stories about legal access inequality — Nyay.AI as the solution narrative
- Submit to ProductHunt India community, YC Hacker News — for early-adopter tech users who will share
- Participate in NASSCOM, iSPIRT events — build B2B credibility for corporate and government pipeline

Phase 3: Paid + Partnership Scale (Month 6–12)

- Meta ads targeting specific legal pain points: 'Landlord not returning deposit?', 'Employer not paying salary?' — direct to WhatsApp bot
- Google Search ads for high-intent queries: 'consumer complaint format India', 'RTI application kaise likhe', 'legal notice format Hindi'
- NBFC and bank partnerships for employee legal helpline: offer as branded benefit to their staff (B2B2C)
- State government legal aid board partnerships for subsidised citizen access

11. Risk Analysis and Mitigation

Risk	Severity	Probability	Mitigation
Legal advice boundary violation (Advocates Act)	HIGH	MEDIUM	Strict system prompt enforcement; 'information not advice' framing; legal review of all template responses; insurance
LLM hallucination producing wrong legal guidance	HIGH	LOW (with RAG)	Temperature=0; citation-only output rule; confidence scoring; escalation to lawyer for low-confidence queries; user disclaimer
Government blocks scraping of court portals	MEDIUM	LOW	Use official APIs where available; cache aggressively; mirror corpus independently; build legal data partnerships
Manupatra or SCC Online launches competing citizen product	MEDIUM	LOW	Their DNA is B2B; they have never served citizens; WhatsApp distribution and regional language depth are moats they can't copy fast
Data freshness — user relies on outdated law	HIGH	MEDIUM	Daily scraper jobs for case law; weekly for circulars; 'last updated' timestamps on every response; always cite source date
Unit economics broken at scale (Claude API costs)	MEDIUM	MEDIUM	Route simple queries to cheaper Llama-3 (self-hosted); reserve Claude for complex synthesis; implement response caching for common queries
WhatsApp Policy violation / account suspension	HIGH	LOW	Maintain compliance with Meta's commerce policy; no unsolicited messages; build email + web as backup channels from Day 1

12. Team and Execution Requirements

12.1 Founding Team Requirements

This business sits at the intersection of law, AI engineering, and consumer product. The founding team must cover all three. A typical founding team profile:

Technical Co-founder (CTO)	Strong Python/ML background; experience with LLMs, RAG, and production ML pipelines. Prior experience with large-scale data pipelines is a plus. The RAG system is the core product — this person builds and owns it.
Product/Growth Co-founder (CEO)	Consumer product experience; ideally prior startup or high-growth company experience. Must understand Bharat users (not just metro users). Hindi fluency is a strong plus for understanding user needs.
Legal Advisor (Part-time / Advisor)	Practicing advocate (Supreme Court or High Court level) who can review template responses for legal accuracy; guide the 'information vs advice' boundary; build credibility with bar associations. Equity advisor role initially.

12.2 First Hires (Post-Seed)

- Data Engineer: builds and maintains scraper pipelines and data ingestion; Python + Scrapy + Airflow experience
- Full-stack Engineer: Next.js frontend + FastAPI backend; WhatsApp integration experience
- Hindi Content Lead: reviews AI outputs for Hindi accuracy; creates Hindi-language training examples for prompt tuning
- Legal Content Reviewer: junior advocate or law graduate; reviews RAG outputs for accuracy and boundary compliance
- Growth Lead: consumer growth with B2C SaaS experience; WhatsApp community building, regional language social media

12.3 Funding Requirements

Round	Target Amount	Use of Funds
Pre-Seed / Bootstrapped	₹0–50 Lakh	Founders build MVP; cover infrastructure costs; first 1,000 users
Seed	₹2–4 Crore	First 4 hires; 12 months runway; corpus expansion; marketing for 1 lakh users; B2G sales pipeline
Series A	₹20–40 Crore	Full team of 25–30; national WhatsApp marketing; lawyer marketplace build-out; B2B product; 10M user target

13. Conclusion — Why Now, Why This

Three forces have converged in 2026 to make Nyay.AI viable in a way it simply was not three years ago:

- **Force 1:** Indic language AI has crossed the quality threshold. Sarvam AI's models understand legal Hindi, Tamil, and Telugu with the accuracy needed for reliable guidance. This was not true in 2022.
- **Force 2:** WhatsApp Business API has matured. The infrastructure to deliver AI at scale to 500 million Indian users without requiring them to install an app or visit a website now exists, is stable, and is affordable.
- **Force 3:** India's legal data is fully digitised and publicly accessible. Every major court now publishes judgments online. India Code has all legislation as machine-readable XML. Government scheme data is on a unified portal. The corpus needed to build a credible RAG system exists and is free.

The competitive landscape has not responded. Manupatra spent 25 years building for lawyers and is now layering AI on top of that B2B product. SCC Online is a publisher adding AI features. CaseMine is a research tool. None of them is building for the person on the other side of the lawyer-client table — the 1.4 billion Indians who are those clients.

The last word

In India, access to justice is a constitutional right under Article 39A. In practice, it is a privilege available to those who can afford lawyers, speak English, and live in cities with courts. Nyay.AI doesn't replace lawyers or courts — it closes the information gap that prevents 80% of Indians from even knowing what rights they have. That is a problem worth solving. It also happens to be a very large business.

— END OF DOCUMENT —